Course notes

Sums and transformations

For a finite random variable X, taking the values x_i with probabilities p_i,

$$E(X) = \sum_i x_i p_i, \qquad V(X) = \sum_i (x_i - E(X))^2 p_i = E(X^2) - E(X)^2.$$

The standard deviation is $\sigma(X) = \sqrt{V(X)}$. Expanding the square in $V(aX+b)$, the shift b cancels out and the factor a is squared. For finite random variables X, Y and reals a, b,

$$E(X+Y) = E(X) + E(Y), \quad E(aX+b) = aE(X) + b, \quad V(aX+b) = a^2 V(X).$$

Linearity of expectation does not require independence. If X, Y are independent, $V(X+Y) = V(X) + V(Y)$ (admitted). Without this assumption, the formula cannot be applied. The standard deviation of $aX+b$ is $|a|\sigma(X)$.

Example of dependence. If $Y = X$, then $V(X+Y) = V(2X) = 4V(X)$, and not 2 $V(X)$ as soon as $V(X) \neq 0$.

Proof for the binomial distribution. Write $X = B_1 + \cdots + B_n$ with the independent indicator variables of successes. Each B_i satisfies $E(B_i) = p$, $E(B_i^2) = p$ and $V(B_i) = p - p^2$. By linearity and additivity under independence, $E(X) = np$, $V(X) = np(1-p)$.

Sample and mean

A sample of size n of a distribution is a family of n independent variables with the same distribution. If they have expectation μ and variance σ^2, let $S_n = \sum_{i=1}^{n} X_i$ & $M_n = \frac{S_n}{n}$.

Sum and mean

$$E(S_n) = n\mu, \quad V(S_n) = n\sigma^2, \qquad E(M_n) = \mu, \quad V(M_n) = \frac{\sigma^2}{n}, \quad \sigma(M_n) = \frac{\sigma}{\sqrt{n}}.$$

To divide the standard deviation of the mean by 2, multiply the sample size by 4. The variance of the sum increases; that of the mean decreases.

Concentration

Bienaymé-Chebyshev inequality

For a variable with expectation μ, variance σ^2, and $\delta > 0$,

$$P(|X - \mu| \geq \delta) \leq \frac{\sigma^2}{\delta^2}.$$

A bound greater than 1 remains true but gives no further information; we can replace it with 1.

Proof in the finite setting. In $V(X) = \sum_x (x-\mu)^2 P(X = x)$, keep only the values satisfying $|x - \mu| \geq \delta$. Each square is then at least δ^2. Hence $V(X) \geq \delta^2 P(|X - \mu| \geq \delta)$.

Law of large numbers

Applied to the mean of a sample,

$$P(|M_n - \mu| \geq \delta) \leq \frac{\sigma^2}{n\delta^2} \xrightarrow[n\to\infty]{} 0.$$

This is the law of large numbers in the setting considered: for a fixed precision, the probability of a deviation at least equal to this precision tends to zero. This result does not say that each realisation approaches μ monotonically.

Choosing a sufficient sample size

To ensure $P(|M_n - \mu| < \delta) \geq 1 - \alpha$, it suffices to choose

$$n \geq \frac{\sigma^2}{\alpha\delta^2}.$$

Take the ceiling. It is the smallest integer satisfying this bound , not necessarily the minimum size for the exact probability.

For a success frequency, $\mu = p$ and $\sigma^2 = p(1-p) \leq \frac{1}{4}$. If p is unknown, $n \geq \frac{1}{4\alpha\delta^2}$ is a uniform guarantee.

From the individual distribution to the mean

14.1 Describing a non-binomial variable

First use: S62.

Worked example · A repair cost

A cost X in euros takes the values 0, 20 and 50 with probabilities 0,5, 0,3 and 0,2. $E(X) = 16$, $E(X^2) = 620$, therefore $V(X) = 620 - 16^2 = 364$ in euros squared. The std dev is $\sqrt{364} \simeq 19{,}08$ euros. The mean 16 is not one of the possible values of X; it describes a theoretical average over many realisations.

A1 · Immediate application X takes values 1, 2, 4 with probabilities $1/2, 1/4, 1/4$. Check the distribution and calculate the expectation, variance and standard deviation.

14.2 A discount also changes the spread

First use: S62.

Property · Affine transformation

For $Y = aX + b$, $E(Y) = aE(X) + b$ and $V(Y) = a^2V(X)$. The standard deviation is $|a|\sigma(X)$. Adding a constant shifts all values without changing their spread.

Worked example · Cost after a discount

For the previous cost, a 20% discount followed by 5 euros of fees gives $Y = 0{,}8X + 5$. Thus $E(Y) = 17.8$ and $V(Y) = 0{,}64 \times 364 = 232.96$. The variance is neither $0.8V(X) + 5$ nor $0.8\,V(X)$.

A2 · Partial guidance For the distribution in the previous exercise, calculate E and V of $Y = 3X - 2$, then of $Z = -X$. Explain why multiplying by -1 does not change the variance.

14.3 Independence matters for the variance

First use: S64.

Property · Sum and mean

For variables $X_1, \dots, X_n$ with the same expectation μ and the same variance σ^2, the mean M_n a has expectation μ. If they are independent, $V(M_n) = \sigma^2/n$ and $\sigma(M_n) = \sigma/\overline{n}$. Linearity of the expectation does not require independence.

Proof. With $S_n = \sum X_i$, linearity gives $E(S_n) = n\mu$. If indep., the admitted additivity of variances gives $V(S_n) = n\sigma^2$. Multiplying by $1/n$ for the mean multiplies the expectation by $1/n$ and the variance by $1/n^2$. For a Bernoulli variable, E=p and $V = p(1-p)$; the binomial variable, sum of n independent Bernoulli variables, therefore has E=np and $V = np(1-p)$. □

Worked example · Mean of a hundred costs

For a hundred independent costs with the same distribution, expectation 16 and variance 364, the total has expectation 1600 and variance 36400. The average cost has expectation 16 and variance 3.64, hence a standard deviation of about 1.91 euros.

A3 · Counterexample X takes the value 0 or 1 with probability 1 /2 each. Compare the variances of $(X_1 + \cdots + X_n)/n$ for independent copies, then for $X_1 = \cdots = X_n = X$.

14.4 A sufficient guarantee, not an exact probability

First use: S73.

Property · Bienaymé-Chebyshev

For $\delta > 0$, $P(|X - \mu| \geq \delta) \leq V(X)/\delta^2$. For a mean of independent variables, $P(|M_n - \mu| \geq \delta) \leq \sigma^2/(n\delta^2)$. Any bound greater than 1 may be replaced by 1.

Worked example · Size guaranteed by the bound

For a success frequency with probability 1 /2, $V(M_n) = 1/(4n)$. We want a gap strictly less than 0 ,1 with probability at least 0 ,95. It suffices to have $25/n \leq 0{,}05$, so $n \geq 500$. This gives the smallest integer satisfying this sufficient inequality; it does not prove that 499 observations actually fail.

A4 · Immediate application A variable has variance 9. What independent sample size guarantees, by this bound, a deviation of the mean strictly less than 0 ,5 with a probability at least 0.99?

A5 · Compare two pieces of information For $X \sim \mathcal{B}(10; 1/2)$ and $M = X/10$, calculate exactly $P(|M - 1/2| \geq 0{,}3)$ and compare with the Chebyshev bound. Work with the exact integer bounds.

14.5 Several samples, several means

First use: S74 and S77.

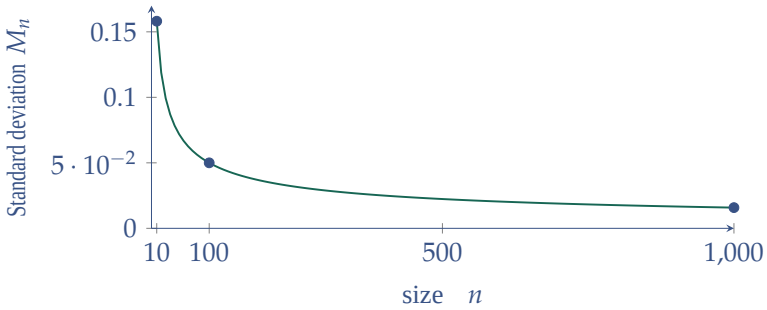


Figure 14.1. Theoretical curve 1/ $(2\sqrt{n})$: multiplying n by 4 divides the standard deviation by 2, not by 4.

Worked example · What is being compared

With n=100, the theoretical standard deviation of a success frequency with probability 1/2 is 0.05. A series of 1,000 experiments of 100 trials produces 1,000 means. It is the standard deviation of this series of means that is compared with 0.05; the 100,000 individual results have a different dispersion.

A6 · Practical work and interpretation With the given program, simulate 1,000 independent samples of sizes 10, 100 and 1,000 for a Bernoulli variable with parameter 1 /2. For each size, compare the observed standard deviation of the means with 1 /(2 $\sqrt{n}$) and count deviations of at least 0 ,1. Why does a single trajectory of frequencies not answer this question?

A7 · Refute "The law of large numbers states that the frequency gets closer to the expectation at each new draw." Give a short sequence of results that temporarily moves the frequency away from 1/2, then restate the result correctly.

One minute, spoken aloud. Write a sentence containing "precision", "risk" and "sufficient size", then distinguish this statement from a certainty for a particular sample.

Delayed review · S90. In a synthesis problem, first identify the distribution or the variance of the mean, before choosing the threshold.

Investigation

Fluctuations of a frequency - 40 min.

A fair coin is tossed independently. For every integer $n \geq 1$, let F_n the frequency of " heads" among the n first tosses.

1. Write a program that simulates 1000 tosses, then plot the points $(n; F_n)$ for $1 \leq n \leq 1000$. Is gap to $\frac{1}{2}$ necessarily decreasing at each toss?
2. Express $E(F_n)$ and $V(F_n)$ as a function of n.
3. Using the Bienaymé-Chebyshev inequality, bound $P(|F_{1000} - \frac{1}{2}| \geq 0{,}05)$.
4. Determine an integer n sufficient to ensure $P(|F_n - \frac{1}{2}| < 0{,}05) \geq 0{,}99$.
5. State what the simulated trajectory shows and what the inequality obtained guarantees. Does a single simulation constitute a proof of this guarantee?